

Short Course on Deep Learning and Convolutional Neural Network

Day 1: Machine Learning Basics
Introduction to Image Processing

Type of Machine Learning Systems

Type of Machine Learning Systems

Supervised Learning
Unsupervised Learning
Semi-supervised Learning
Reinforcement Learning

Depending on whether the system is trained with human supervision

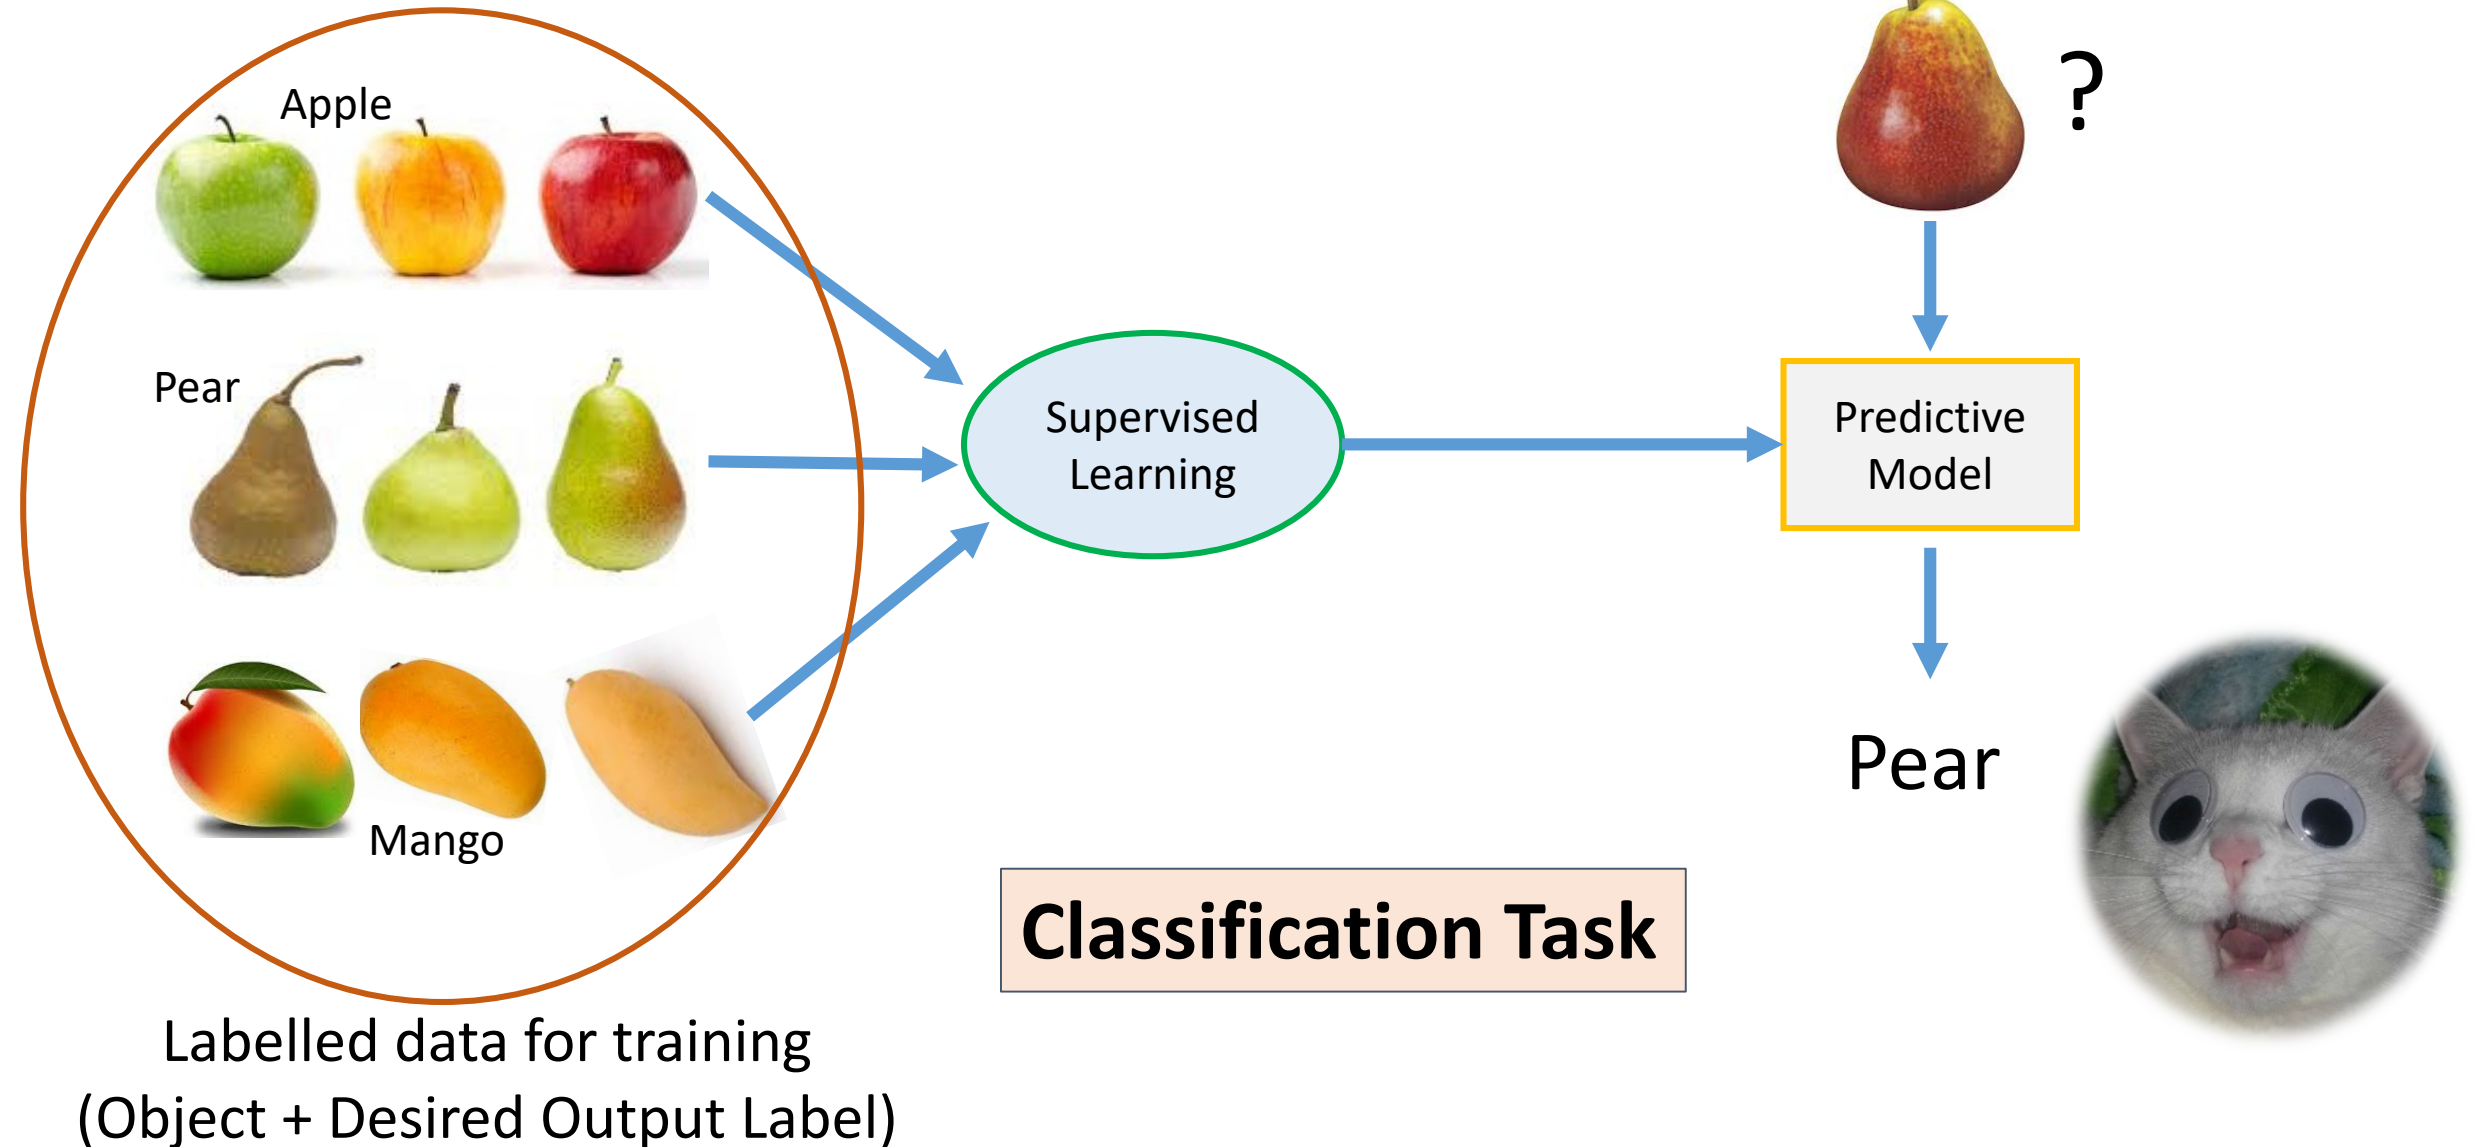
Whether System can learn on the fly

Batch and Online Learning

Instance-based and Model-based Learning

Comparing data points or detect patterns in training data to build a predictive model

Supervised Learning



Supervised Learning

House Price prediction

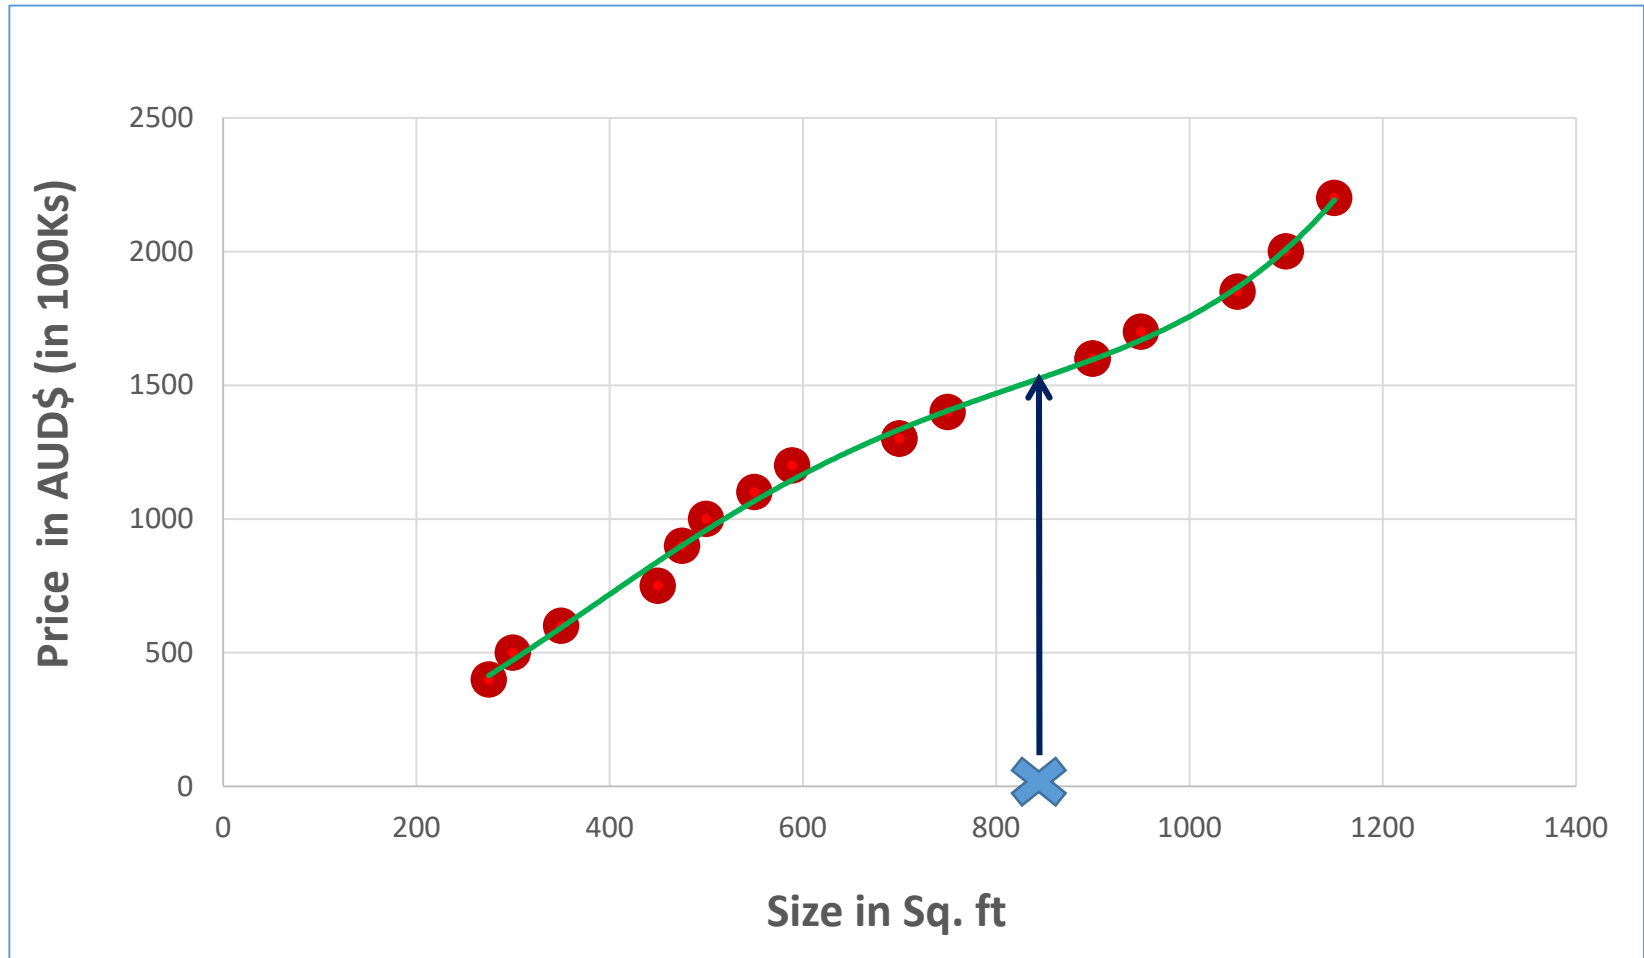
Feature:

- Size of the house

To Predict:

- Price of the house

Regression Task

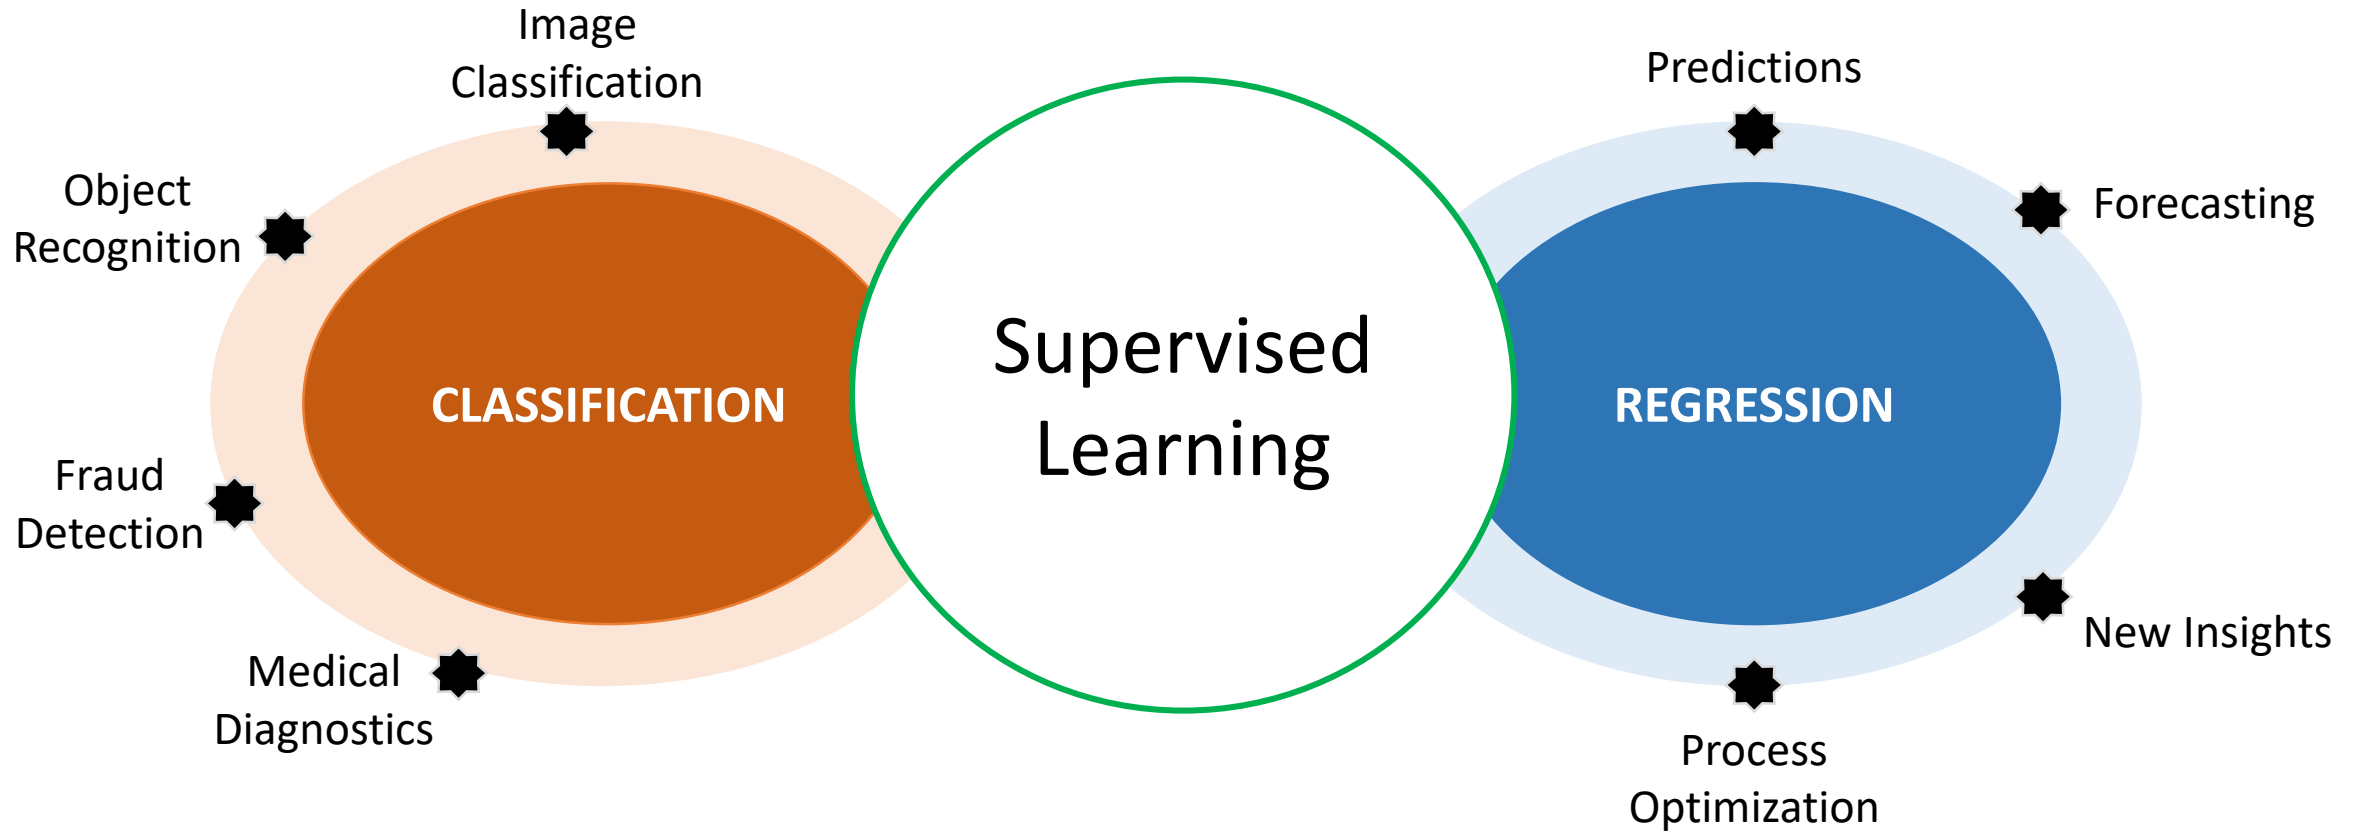


Supervised Learning

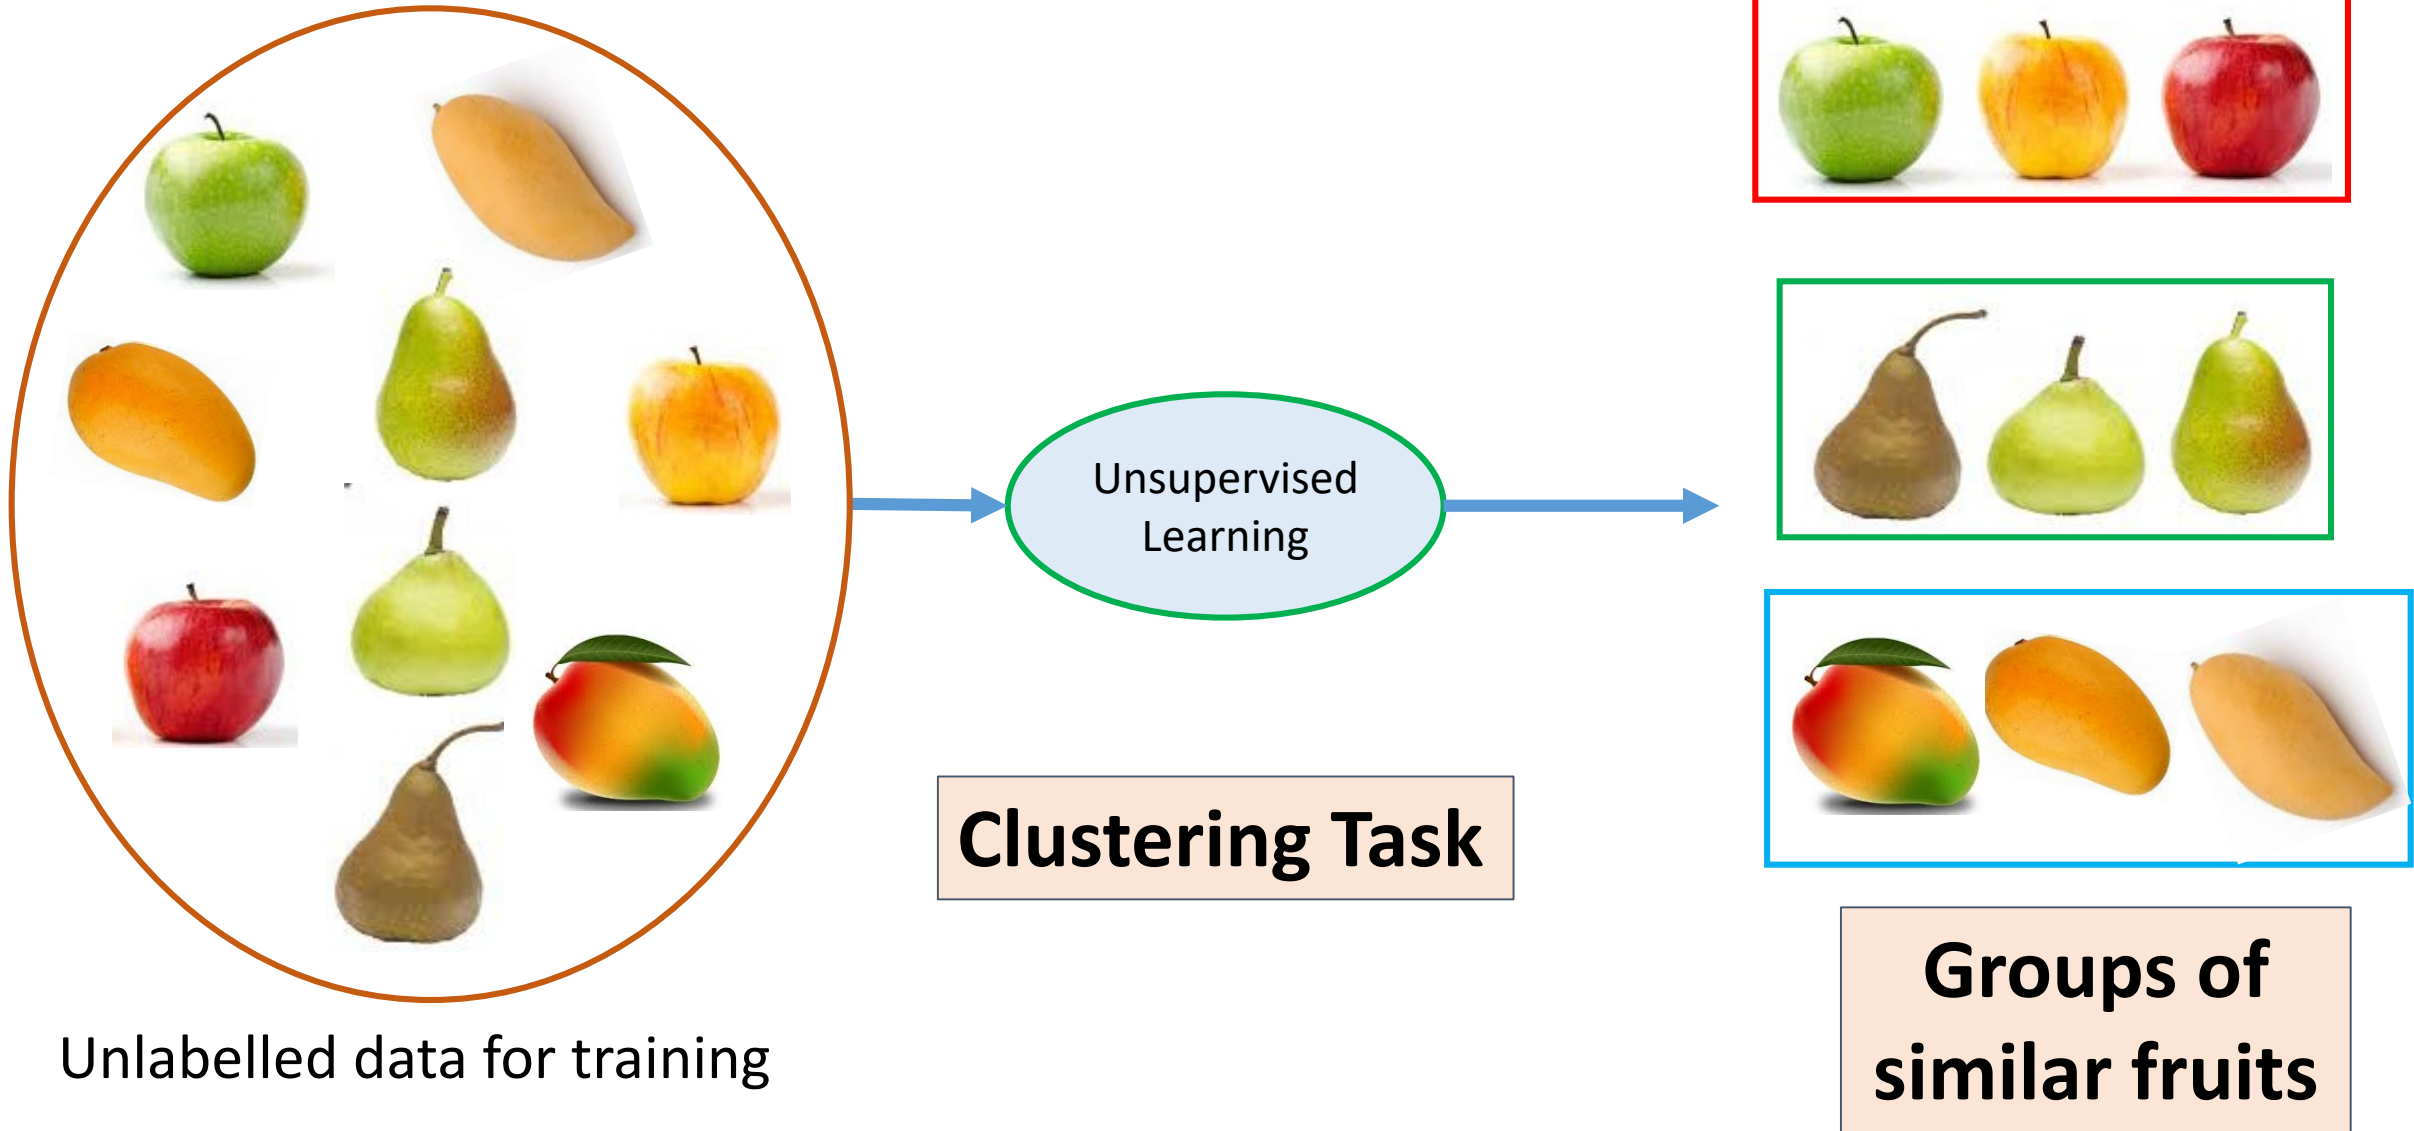
Important Algorithms:

- K-Nearest Neighbours
- Logistic regression
- Support Vector Machines (SVMs)
- Neural Networks (*some of them can be unsupervised)

Supervised Learning Examples



Unsupervised Learning



Unsupervised Learning

Important Algorithms:

- k-means
- Expectation Maximization

Support Vector Machines (SVM)

- A Support Vector Machine is a very powerful and versatile Machine Learning model, capable of performing linear or non-linear classification, regression, and also outlier detection.
- Defined by a separating hyperplane
- Suitable for small or medium sized datasets

Reference and Pre-Reading:

Theory: <https://medium.com/machine-learning-101/chapter-2-svm-support-vector-machine-theory-f0812effc72>

Implementation: <https://medium.com/machine-learning-101/chapter-2-svm-support-vector-machine-coding-edd8f1cf8f2d>

Support Vector Machines (SVM)

Example: Using sklearn for SVM classification (Partial code snippet)

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm, datasets

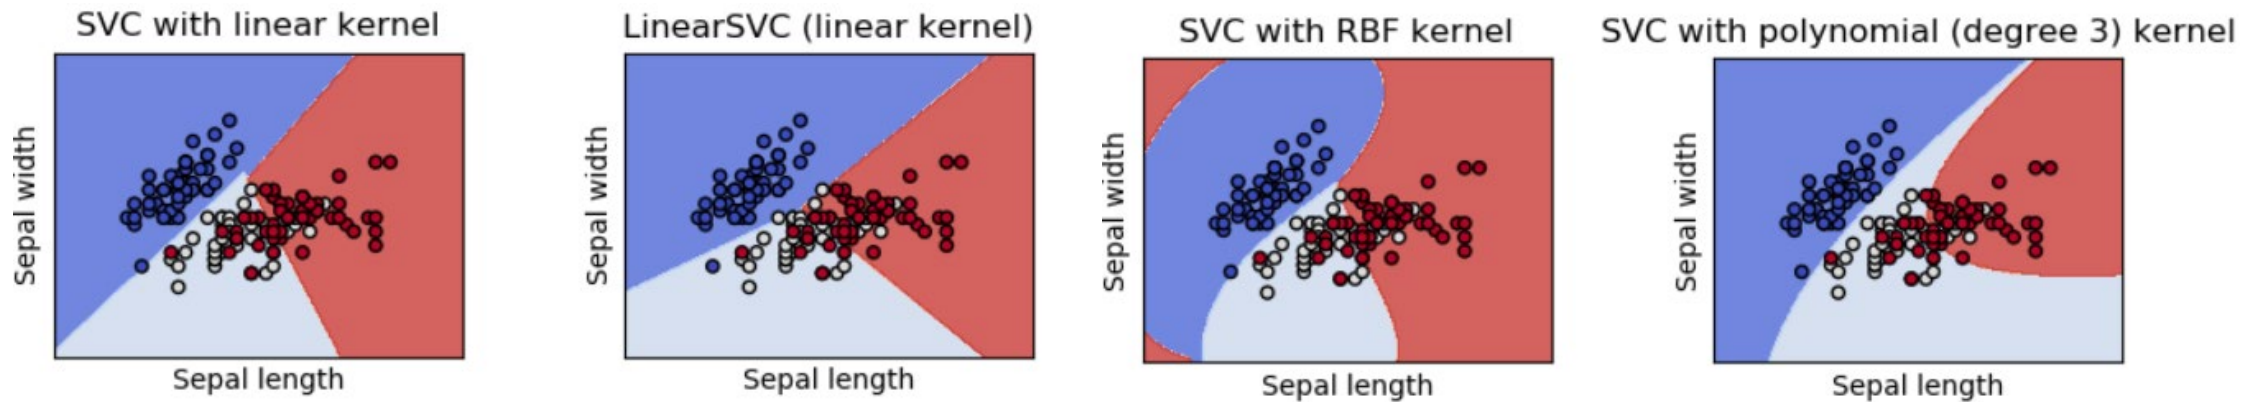
# import some data to play with
iris = datasets.load_iris()
# Take the first two features. We could avoid this by using a two-dim dataset
X = iris.data[:, :2]
y = iris.target

# we create an instance of SVM and fit out data. We do not scale our
# data since we want to plot the support vectors
C = 1.0 # SVM regularization parameter
models = (svm.SVC(kernel='linear', C=C),
          svm.LinearSVC(C=C),
          svm.SVC(kernel='rbf', gamma=0.7, C=C),
          svm.SVC(kernel='poly', degree=3, C=C))
models = (clf.fit(X, y) for clf in models)
```

Reference: https://scikit-learn.org/stable/auto_examples/svm/plot_iris.html#sphx-glr-auto-examples-svm-plot-iris-py
https://en.wikipedia.org/wiki/Iris_flower_data_set

Support Vector Machines (SVM)

Example: Using sklearn for SVM classification



Reference: https://scikit-learn.org/stable/auto_examples/svm/plot_iris.html#sphx-glr-auto-examples-svm-plot-iris-py
https://en.wikipedia.org/wiki/Iris_flower_data_set

K-Nearest Neighbour (KNN)


- A simple supervised learning algorithm.
- Can be used for both classification and regression
- Non-parametric: doesn't make any assumption on the data distribution
- Training data is retained to make future predictions

K-Nearest Neighbour (KNN)

How does it work?

- Computes distance between the new sample and all training samples
- Distance measure: Euclidean, Manhattan etc.
- Picks 'k' entries in the training set which are closest to the new sample
- Majority voting decides the class of the new sample

K-Nearest Neighbour (KNN)



Evaluation Metrics

- **Precision & Recall**

What are the “correct” cells?

- *TN*: (Number of True Negatives),
i.e., patients who did *not* have cancer whom we correctly diagnosed as *not* having cancer.
- *TP*: (Number of True Positives),
i.e., patients who did have cancer whom we correctly diagnosed as having cancer

Precision: TP/Cancer Diagnoses

		Diagnoses	
		No Cancer	Cancer
True state	No Cancer	TN	FP
	Cancer	FN	TP

Recall: TP/Cancer True States

Evaluation Metrics

- **Precision & Recall**

what are the “error” cells are:

- *FN*: (Number of False Negatives),
i.e., patients who did have cancer whom we incorrectly diagnosed as *not* having cancer
- *FP*: (Number of False Positives),
i.e., patients who did *not* have cancer whom we incorrectly diagnosed as having cancer

Precision: $TP / \text{Cancer Diagnoses}$

		Diagnoses	
		No Cancer	Cancer
True state	No Cancer	TN	FP
	Cancer	FN	TP

Recall: $TP / \text{Cancer True States}$


$$\text{Precision} = (TP) / (TP + FP)$$

$$\text{Recall} = (TP) / (TP + FN)$$

Evaluation Metrics

- **Intersection over Union (IoU):**

Intersection over Union is a metric used for the evaluation of an object detector, i.e. how good is the predicted bounding box for an object detected closely matches

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$





Image Processing Basics

Image Processing Basics

What is a digital image?

- Digital images are made of picture elements called ***Pixels***.
- It is an array, or a matrix of ***Pixels*** arranged in *columns* and *rows*.
- Each ***Pixel*** has its own *intensity* value, or *brightness*
- Intensity values in digital images are defined by ***bits***
- For a standard 8 bits image, a pixel can have $2^8 = 256$ (0 – 255) values.
- Black & White images have a single 8-bits intensity range.

Image Processing Basics

Image dimension = 5 X 5

$f(2, 3) = 170$ (Pixel/intensity value)

Hence, an image may be defined as a 2D function $f(x, y)$, where, x and y are spatial co-ordinates, and the amplitude of f at (x, y) is the intensity or Gray level of the image at that point/pixel.

170	170	55	170	170
170	55	170	55	170
170	55	170	55	170
55	140	140	140	55
55	170	170	170	55

5 X 5 Gray scale image (8 bit)

Image Processing Basics

Image dimension = 5 X 5 X 3

No. of Channels = 3

Since, RGB image contains 3 X 8-bits of intensities, they are referred to as 24-bit colour images.

So, 24-bit colour depth

= 8 X 8 X 8 bits

= 256 X 256 X 256 colours

= ~16 million colours

Blue	170	170	55	170	170			
Green	170	170	55	170	170	70		
Red	170	170	55	170	170	70	70	
	170	55	170	55	170	70	55	
	170	55	170	55	170	55	55	
	55	140	140	140	55	55		
	55	170	170	170	55			

5 X 5 X 3 colour image (24 bit)


Image Thresholding

- Easiest method for image segmentation!
- Converts gray-scale image into a binary image
If $f(x,y) > \text{Threshold}$, then $f(x,y) = 0$ else $f(x,y) = 255$

Binary Image (8-bit) has only two possible values of pixel intensity (0 and 1, or B & W)

170	170	55	170	170
170	55	170	55	170
170	55	170	55	170
55	140	140	140	55
55	170	170	170	55

Threshold = 100



255	255	0	255	255
255	0	255	0	255
255	0	255	0	255
0	255	255	255	0
0	255	255	255	0





Image Thresholding



Original Image

Thresholding



Binary Image

Image Thresholding methods

- **Histogram shape:**


Peaks, valleys and curvature of the histogram are analysed.

- **Clustering based:**

The ¹Otsu method, very good for bimodal distribution

- **Adaptive thresholding:**


Instead of a single threshold, have thresholds for different regions in the image



Edge Detection

What is an edge?

- The points/pixels in an image where brightness/intensities changes sharply
- A simple and fundamental tools in image processing and computer vision, useful in feature detection/extraction



How to find Edges in images? (Convolution)

100	100	100	0	0	0
100	100	100	0	0	0
100	100	100	0	0	0
100	100	100	0	0	0
100	100	100	0	0	0
100	100	100	0	0	0

6 X 6 dimension image

*

Convolution
operator

1	0	-1
1	0	-1
1	0	-1


3 X 3 filter/Kernel



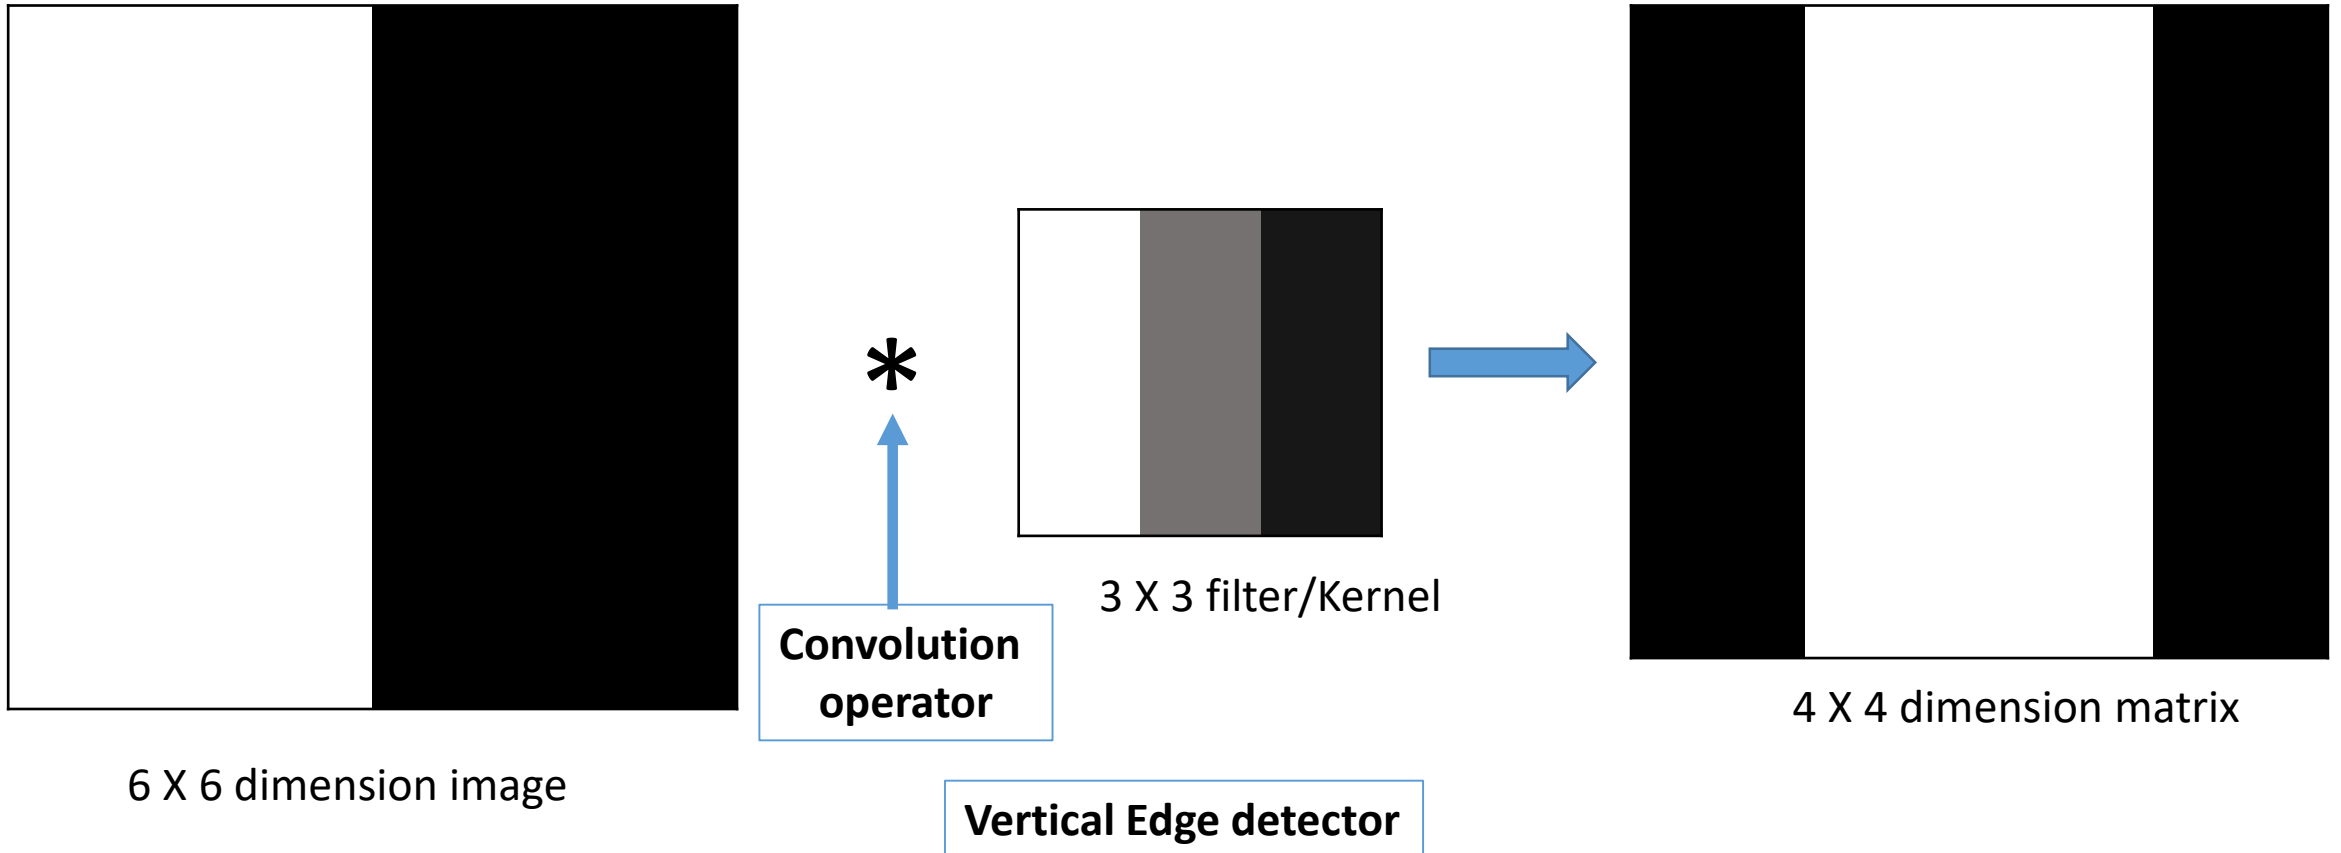
0	300	300	0
0	300	300	0
0	300	300	0
0	300	300	0

4 X 4 dimension matrix

How to find Edges in images? (Convolution)



How to find Edges in images? (Convolution)



How to find Edges in images? (Convolution)

100	100	100	100	100	100
100	100	100	100	100	100
100	100	100	100	100	100
0	0	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

6 X 6 dimension image

*

Convolution
operator

1	1	1
0	0	0
-1	-1	-1

3 X 3 filter/Kernel




0	0	0	0
300	300	300	300
300	300	300	300
0	0	0	0

4 X 4 dimension matrix

Horizontal Edge detector

How to find Edges in images? (Convolution)



Sobel edge detection

1	0	-1
1	0	-1
1	0	-1

3 X 3 filter/Kernel
For Vertical edges

1	1	1
0	0	0
-1	-1	-1

3 X 3 filter/Kernel
For Horizontal edges

Prewitt Filters

1	2	1
0	0	0
-1	-2	-1

3 X 3 filter/Kernel
For Horizontal edges

1	0	-1
2	0	-2
1	0	-1

3 X 3 filter/Kernel
For Vertical edges


Sobel Filters

Sobel edge detection


Original



Sobel X



Sobel Y



Convolutions in CNN

- Convolutions are very important operation in a Convolutional Neural Networks (CNN)
- Filters weights are not fixed, but learned during the training operations of a CNN for a specific task!
- Multiple filters are used in CNNs